

We use cookies on kaggle to deliver our services, analyze web traffic, and improve your experience on the site. By using kaggle, you agree to our use of cookies.

Got it

Learn more

kaggle

Search

Competitions

Datasets


Kernels


Discussion

Learn

...

Sign In





Text Classification with QRNN

Python notebook using data from [Toxic Comment Classification Challenge](#) · 574 views ·  
🔖 deep learning, text data, text mining

^

1

Fork

4

...

Version 4

[4 commits](#)

forked from Text Classification with Atrous Convolutions

Submission	Private Score	Public Score
✓ Ran successfully Submitted by Kevin Mader 9 months ago	0.50000	0.50000

Notebook

Data

Output

Log

Comments

## Overview

Here we show how well a QRNN works when compared with standard CNN and LSTM models. The parent is a 1D Convolution and other models show the stacked LSTM performance

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from keras.models import Model
from keras.layers import Dense, Embedding, Input
from keras.layers import Conv1D, GlobalMaxPool1D, Dropout, concatenate
from keras.preprocessing import text, sequence
from keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
/opt/conda/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
  from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

## QRNN Implementation

We use a reference implementation of the QRNN layer from [https://github.com/DingKe/nn\\_playground/blob/master/q rnn/q rnn.py](https://github.com/DingKe/nn_playground/blob/master/q rnn/q rnn.py) ([https://github.com/DingKe/nn\\_playground/blob/master/q rnn/q rnn.py](https://github.com/DingKe/nn_playground/blob/master/q rnn/q rnn.py))

```
In [2]: from keras import backend as K
from keras import activations, initializers, regularizers, constraints
from keras.layers import Layer, InputSpec
from keras.utils.conv_utils import conv_output_length

def _dropout(x, level, noise_shape=None, seed=None):
    x = K.dropout(x, level, noise_shape, seed)
    x *= (1. - level) # compensate for the scaling by the dropout
    return x

class QRNN(Layer):
    '''Quasi RNN
    # Arguments
        units: dimension of the internal projections and the fir
```

al output.

# References

- [Quasi-recurrent Neural Networks](http://arxiv.org/abs/1611.01576)

'''

```
def __init__(self, units, window_size=2, stride=1,
              return_sequences=False, go_backwards=False,
              stateful=False, unroll=False, activation='tan
h',
              kernel_initializer='uniform', bias_initialize
='zero',
              kernel_regularizer=None, bias_regularizer=None
,
              activity_regularizer=None,
              kernel_constraint=None, bias_constraint=None,
              dropout=0, use_bias=True, input_dim=None, inp
t_length=None,
              **kwargs):
    self.return_sequences = return_sequences
    self.go_backwards = go_backwards
    self.stateful = stateful
    self.unroll = unroll

    self.units = units
    self.window_size = window_size
    self.strides = (stride, 1)

    self.use_bias = use_bias
    self.activation = activations.get(activation)
    self.kernel_initializer = initializers.get(kernel_init
alizer)
    self.bias_initializer = initializers.get(bias_initiali
er)
    self.kernel_regularizer = regularizers.get(kernel_regu
arizer)
    self.bias_regularizer = regularizers.get(bias_regulari
er)
    self.activity_regularizer = regularizers.get(activity_
egularizer)
    self.kernel_constraint = constraints.get(kernel_constr
int)
    self.bias_constraint = constraints.get(bias_constraint

    self.dropout = dropout
    self.supports_masking = True
    self.input_spec = [InputSpec(ndim=3)]
    self.input_dim = input_dim
    self.input_length = input_length
    if self.input_dim:
        kwargs['input_shape'] = (self.input_length, self.in
put_dim)
    super(QRNN, self).__init__(**kwargs)

    def build(self, input_shape):
```

```

    if isinstance(input_shape, list):
        input_shape = input_shape[0]

    batch_size = input_shape[0] if self.stateful else None
    self.input_dim = input_shape[2]
    self.input_spec = InputSpec(shape=(batch_size, None, self.input_dim))
    self.state_spec = InputSpec(shape=(batch_size, self.units))

    self.states = [None]
    if self.stateful:
        self.reset_states()

    kernel_shape = (self.window_size, 1, self.input_dim, self.units * 3)
    self.kernel = self.add_weight(name='kernel',
                                   shape=kernel_shape,
                                   initializer=self.kernel_initializer,
                                   regularizer=self.kernel_regularizer,
                                   constraint=self.kernel_constraint)

    if self.use_bias:
        self.bias = self.add_weight(name='bias',
                                     shape=(self.units * 3),
                                     initializer=self.bias_initializer,
                                     regularizer=self.bias_regularizer,
                                     constraint=self.bias_constraint)

    self.built = True

    def compute_output_shape(self, input_shape):
        if isinstance(input_shape, list):
            input_shape = input_shape[0]

        length = input_shape[1]
        if length:
            length = conv_output_length(length + self.window_size - 1,
                                         self.window_size, 'valid',
                                         self.strides[0])

        if self.return_sequences:
            return (input_shape[0], length, self.units)
        else:
            return (input_shape[0], self.units)

    def compute_mask(self, inputs, mask):

```

```

    if self.return_sequences:
        return mask
    else:
        return None

def get_initial_states(self, inputs):
    # build an all-zero tensor of shape (samples, units)
    initial_state = K.zeros_like(inputs) # (samples, time
    steps, input_dim)
    initial_state = K.sum(initial_state, axis=(1, 2)) # (s
    amples,)
    initial_state = K.expand_dims(initial_state) # (sample
    s, 1)
    initial_state = K.tile(initial_state, [1, self.units])
    # (samples, units)
    initial_states = [initial_state for _ in range(len(self
    .states))]
    return initial_states

def reset_states(self, states=None):
    if not self.stateful:
        raise AttributeError('Layer must be stateful.')
    if not self.input_spec:
        raise RuntimeError('Layer has never been called '
                           'and thus has no states.')

    batch_size = self.input_spec.shape[0]
    if not batch_size:
        raise ValueError('If a QRNN is stateful, it needs
    o know '
                           'its batch size. Specify the batch
    size '
                           'of your input tensors: \n'
                           '- If using a Sequential model, '
                           'specify the batch size by passing
    ,
                           'a `batch_input_shape` '
                           'argument to your first layer.\n'
                           '- If using the functional API, sp
    ecify '
                           'the time dimension by passing a
                           ``batch_shape` argument to your I
    put layer.')

    if self.states[0] is None:
        self.states = [K.zeros((batch_size, self.units))
                        for _ in self.states]
    elif states is None:
        for state in self.states:
            K.set_value(state, np.zeros((batch_size, self.u
    nits)))
    else:
        if not isinstance(states, (list, tuple)):
            states = [states]

```

```

        if len(states) != len(self.states):
            raise ValueError('Layer ' + self.name + ' expected ' +
                               str(len(self.states)) + ' states, '
                               'but it received ' + str(len(states)) +
                               'state values. Input received ' +
                               str(states))
        for index, (value, state) in enumerate(zip(states, self.states)):
            if value.shape != (batch_size, self.units):
                raise ValueError('State ' + str(index) +
                                   ' is incompatible with layer ' +
                                   self.name + ': expected shape=' +
                                   str((batch_size, self.units)) +
                                   ', found shape=' + str(value.shape))
            K.set_value(state, value)

    def __call__(self, inputs, initial_state=None, **kwargs):
        # If `initial_state` is specified,
        # and if it a Keras tensor,
        # then add it to the inputs and temporarily
        # modify the input spec to include the state.
        if initial_state is not None:
            if hasattr(initial_state, '_keras_history'):
                # Compute the full input spec, including state
                input_spec = self.input_spec
                state_spec = self.state_spec
                if not isinstance(state_spec, list):
                    state_spec = [state_spec]
                self.input_spec = [input_spec] + state_spec

                # Compute the full inputs, including state
                if not isinstance(initial_state, (list, tuple)):
                    initial_state = [initial_state]
                inputs = [inputs] + list(initial_state)

            # Perform the call
            output = super(QRNN, self).__call__(inputs, **kwargs)

            # Restore original input spec
            self.input_spec = input_spec
            return output
        else:
            kwargs['initial_state'] = initial_state
            return super(QRNN, self).__call__(inputs, **kwargs)

```

```

def call(self, inputs, mask=None, initial_state=None, training=None):
    # input shape: `(samples, time (padded with zeros), input_dim)`
    # note that the .build() method of subclasses MUST define
    # self.input_spec and self.state_spec with complete input shapes.
    if isinstance(inputs, list):
        initial_states = inputs[1:]
        inputs = inputs[0]
    elif initial_state is not None:
        pass
    elif self.stateful:
        initial_states = self.states
    else:
        initial_states = self.get_initial_states(inputs)

    if len(initial_states) != len(self.states):
        raise ValueError('Layer has ' + str(len(self.states)) +
            ' states but was passed ' +
            str(len(initial_states)) +
            ' initial states.')
    input_shape = K.int_shape(inputs)
    if self.unroll and input_shape[1] is None:
        raise ValueError('Cannot unroll a RNN if the '
            'time dimension is undefined. \n'
            '- If using a Sequential model, '
            'specify the time dimension by passing '
            'an `input_shape` or `batch_input_shape` '
            'argument to your first layer. If '
            'your first layer is an Embedding, you '
            'can also use the `input_length` argument.\n'
            '- If using the functional API, specify '
            'the time dimension by passing a '
            'shape` '
            'or `batch_shape` argument to your '
            'Input layer.')
    constants = self.get_constants(inputs, training=None)
    preprocessed_input = self.preprocess_input(inputs, training=None)

    last_output, outputs, states = K.rnn(self.step, preprocessed_input,
                                           initial_states,
                                           go_backwards=self.go_backwards)

```

```

o_backwards,

                                mask=mask,
                                constants=constant:

,

                                unroll=self.unroll
                                input_length=input.

shape[1])
    if self.stateful:
        updates = []
        for i in range(len(states)):
            updates.append((self.states[i], states[i]))
        self.add_update(updates, inputs)

    # Properly set learning phase
    if 0 < self.dropout < 1:
        last_output._uses_learning_phase = True
        outputs._uses_learning_phase = True

    if self.return_sequences:
        return outputs
    else:
        return last_output

    def preprocess_input(self, inputs, training=None):
        if self.window_size > 1:
            inputs = K.temporal_padding(inputs, (self.window_size-1, 0))
        inputs = K.expand_dims(inputs, 2) # add a dummy dimension

        output = K.conv2d(inputs, self.kernel, strides=self.strides,
                                padding='valid',
                                data_format='channels_last')
        output = K.squeeze(output, 2) # remove the dummy dimension

        if self.use_bias:
            output = K.bias_add(output, self.bias, data_format='channels_last')

        if self.dropout is not None and 0. < self.dropout < 1:
            z = output[:, :, :self.units]
            f = output[:, :, self.units:2 * self.units]
            o = output[:, :, 2 * self.units:]
            f = K.in_train_phase(1 - _dropout(1 - f, self.dropout), f, training=training)
            return K.concatenate([z, f, o], -1)
        else:
            return output

    def step(self, inputs, states):
        prev_output = states[0]

        z = inputs[:, :self.units]

```



```

        f = inputs[:, self.units:2 * self.units]
        o = inputs[:, 2 * self.units:]

        z = self.activation(z)
        f = f if self.dropout is not None and 0. < self.dropout
< 1. else K.sigmoid(f)
        o = K.sigmoid(o)

        output = f * prev_output + (1 - f) * z
        output = o * output

        return output, [output]

def get_constants(self, inputs, training=None):
    return []

def get_config(self):
    config = {'units': self.units,
              'window_size': self.window_size,
              'stride': self.strides[0],
              'return_sequences': self.return_sequences,
              'go_backwards': self.go_backwards,
              'stateful': self.stateful,
              'unroll': self.unroll,
              'use_bias': self.use_bias,
              'dropout': self.dropout,
              'activation': activations.serialize(self.act
vation),
              'kernel_initializer': initializers.serialize
self.kernel_initializer),
              'bias_initializer': initializers.serialize(s
lf.bias_initializer),
              'kernel_regularizer': regularizers.serialize
self.kernel_regularizer),
              'bias_regularizer': regularizers.serialize(s
lf.bias_regularizer),
              'activity_regularizer': regularizers.seriali
e(self.activity_regularizer),
              'kernel_constraint': constraints.serialize(s
lf.kernel_constraint),
              'bias_constraint': constraints.serialize(sel
f.bias_constraint),
              'input_dim': self.input_dim,
              'input_length': self.input_length}
    base_config = super(QRNN, self).get_config()
    return dict(list(base_config.items()) + list(config.it
ms()))

```

In [3]:

```

# define network parameters
max_features = 20000
maxlen = 100

```

## Load and Preprocessing Steps

Here we load the data and fill in the missing values

```
In [4]:
train = pd.read_csv("../input/train.csv")
test = pd.read_csv("../input/test.csv")
train = train.sample(frac=1)

list_sentences_train = train["comment_text"].fillna("Invalid")
values
list_classes = ["toxic", "severe_toxic", "obscene", "threat",
"insult", "identity_hate"]
y = train[list_classes].values
list_sentences_test = test["comment_text"].fillna("Invalid").values
```

## Sequence Generation

Here we take the data and generate sequences from the data

```
In [5]:
tokenizer = text.Tokenizer(num_words=max_features)
tokenizer.fit_on_texts(list(list_sentences_train))
# train data
list_tokenized_train = tokenizer.texts_to_sequences(list_sentences_train)
X_t = sequence.pad_sequences(list_tokenized_train, maxlen=maxlen)
# test data
list_tokenized_test = tokenizer.texts_to_sequences(list_sentences_test)
X_te = sequence.pad_sequences(list_tokenized_test, maxlen=maxlen)
```

```
In [6]:
def build_model(conv_layers = 2, max_dilation_rate = 3):
    embed_size = 128
    inp = Input(shape=(maxlen, ))
    x = Embedding(max_features, embed_size)(inp)
    x = Dropout(0.25)(x)
    x = Conv1D(2*embed_size,
               kernel_size = 3)(x)
    prefilt_x = Conv1D(2*embed_size,
                       kernel_size = 3)(x)
    out_conv = []
    x = prefilt_x
    # strides rate we use here for skip
    for strides in [1, 1, 2]:
        x = QRNN(128*2**(strides),
                  return_sequences = True)
```

```

        return_sequences = True,
        stride = strides,
        dropout = 0.2)(x)
x_f = QRNN(512)(x)
x_b = QRNN(512, go_backwards=True)(x)
x = concatenate([x_f, x_b])
x = Dropout(0.5)(x)
x = Dense(64, activation="relu")(x)
x = Dropout(0.1)(x)
x = Dense(6, activation="sigmoid")(x)
model = Model(inputs=inp, outputs=x)
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['binary_accuracy'])

return model

model = build_model()
model.summary()

```

```

-----
Layer (type)                 Output Shape              Param #
Connected to
-----
input_1 (InputLayer)        (None, 100)               0
-----
embedding_1 (Embedding)     (None, 100, 128)         2560000
input_1[0][0]
-----
dropout_1 (Dropout)         (None, 100, 128)         0
embedding_1[0][0]
-----
conv1d_1 (Conv1D)           (None, 98, 256)          98560
dropout_1[0][0]
-----
conv1d_2 (Conv1D)           (None, 96, 256)          196864
conv1d_1[0][0]
-----
qrnn_1 (QRNN)               (None, 96, 256)          393984
conv1d_2[0][0]
-----
qrnn_2 (QRNN)               (None, 96, 256)          393984
qrnn_1[0][0]
-----
qrnn_3 (QRNN)               (None, 48, 512)          787968

```

```
qrnn_2[0][0]
```

```
-----
qrnn_4 (QRNN)                (None, 512)                1574400
qrnn_3[0][0]
```

```
-----
qrnn_5 (QRNN)                (None, 512)                1574400
qrnn_3[0][0]
```

```
-----
concatenate_1 (Concatenate)   (None, 1024)                0
qrnn_4[0][0]
```

```
qrnn_5[0][0]
```

```
-----
dropout_2 (Dropout)          (None, 1024)                0
concatenate_1[0][0]
```

```
-----
dense_1 (Dense)              (None, 64)                  65600
dropout_2[0][0]
```

```
-----
dropout_3 (Dropout)          (None, 64)                  0
dense_1[0][0]
```

```
-----
dense_2 (Dense)              (None, 6)                   390
dropout_3[0][0]
```

```
=====
Total params: 7,646,150
Trainable params: 7,646,150
Non-trainable params: 0
```

## Train the Model

Here we train the model and use model checkpointing and early stopping to keep only the best version of the model

In [7]:

```
batch_size = 1024 # big beefy GPUs like large batches
epochs = 12
```

```
file_path="weights.hdf5"
checkpoint = ModelCheckpoint(file_path, monitor='val_loss', ve
bose=1, save_best_only=True, mode='min')
```

```
early = EarlyStopping(monitor="val_loss", mode="min", patience=
5)
```

```
callbacks_list = [checkpoint, early] #early
model.fit(X_t, y,
          batch_size=batch_size,
          epochs=epochs,
          shuffle = True,
          validation_split=0.25,
          callbacks=callbacks_list)
```

Train on 119678 samples, validate on 39893 samples

Epoch 1/12

```
119678/119678 [=====] - 244s 2ms/step
- loss: 0.6415 - binary_accuracy: 0.9596 - val_loss: 0.6409 -
val_binary_accuracy: 0.9632
```

Epoch 00001: val\_loss improved from inf to 0.64090, saving model to weights.hdf5

Epoch 2/12

```
119678/119678 [=====] - 232s 2ms/step
- loss: 0.6166 - binary_accuracy: 0.9634 - val_loss: 0.5929 -
val_binary_accuracy: 0.9632
```

Epoch 00002: val\_loss improved from 0.64090 to 0.59286, saving model to weights.hdf5

Epoch 3/12

```
119678/119678 [=====] - 232s 2ms/step
- loss: 0.5711 - binary_accuracy: 0.9634 - val_loss: 0.5498 -
val_binary_accuracy: 0.9632
```

Epoch 00003: val\_loss improved from 0.59286 to 0.54985, saving model to weights.hdf5

Epoch 4/12

```
119678/119678 [=====] - 232s 2ms/step
- loss: 0.5303 - binary_accuracy: 0.9634 - val_loss: 0.5112 -
val_binary_accuracy: 0.9632
```

Epoch 00004: val\_loss improved from 0.54985 to 0.51116, saving model to weights.hdf5

Epoch 5/12

```
119678/119678 [=====] - 232s 2ms/step
- loss: 0.4935 - binary_accuracy: 0.9634 - val_loss: 0.4763 -
val_binary_accuracy: 0.9632
```

Epoch 00005: val\_loss improved from 0.51116 to 0.47631, saving model to weights.hdf5

Epoch 6/12

```
119678/119678 [=====] - 232s 2ms/step
- loss: 0.4604 - binary_accuracy: 0.9634 - val_loss: 0.4449 -
val_binary_accuracy: 0.9632
```

Epoch 00006: val\_loss improved from 0.47631 to 0.44494, saving model to weights.hdf5

```
Epoch 7/12
119678/119678 [=====] - 232s 2ms/step
- loss: 0.4305 - binary_accuracy: 0.9634 - val_loss: 0.4167 -
val_binary_accuracy: 0.9632
```

```
Epoch 00007: val_loss improved from 0.44494 to 0.41669, saving
model to weights.hdf5
```

```
Epoch 8/12
119678/119678 [=====] - 232s 2ms/step
- loss: 0.4037 - binary_accuracy: 0.9634 - val_loss: 0.3912 -
val_binary_accuracy: 0.9632
```

```
Epoch 00008: val_loss improved from 0.41669 to 0.39124, saving
model to weights.hdf5
```

```
Epoch 9/12
119678/119678 [=====] - 232s 2ms/step
- loss: 0.3795 - binary_accuracy: 0.9634 - val_loss: 0.3683 -
val_binary_accuracy: 0.9632
```

```
Epoch 00009: val_loss improved from 0.39124 to 0.36830, saving
model to weights.hdf5
```

```
Epoch 10/12
119678/119678 [=====] - 232s 2ms/step
- loss: 0.3577 - binary_accuracy: 0.9634 - val_loss: 0.3476 -
val_binary_accuracy: 0.9632
```

```
Epoch 00010: val_loss improved from 0.36830 to 0.34764, saving
model to weights.hdf5
```

```
Epoch 11/12
119678/119678 [=====] - 232s 2ms/step
- loss: 0.3380 - binary_accuracy: 0.9634 - val_loss: 0.3290 -
val_binary_accuracy: 0.9632
```

```
Epoch 00011: val_loss improved from 0.34764 to 0.32900, saving
model to weights.hdf5
```

```
Epoch 12/12
119678/119678 [=====] - 232s 2ms/step
- loss: 0.3203 - binary_accuracy: 0.9634 - val_loss: 0.3122 -
val_binary_accuracy: 0.9632
```

```
Epoch 00012: val_loss improved from 0.32900 to 0.31217, saving
model to weights.hdf5
```

Out[7]:

```
<keras.callbacks.History at 0x7f6574d93278>
```

## Test on a few sentences

Show what the model actually predicts for a few test sentences

In [8]:

```
from IPython.display import Markdown, display
```

```

dmd = lambda x: display(Markdown(x))
def show_sentence(sent_idx):
    dmd('# Input Sentence:\n `{}`'.format(list_sentences_train[
sent_idx]))
    c_pred = model.predict(X_t[sent_idx:sent_idx+1])[0]
    dmd('## Positive Categories')
    for k, v, p in zip(list_classes, y[sent_idx], c_pred):
        if v>0:
            dmd('- {}, Prediction: {:.2f}%'.format(k, 100*v,
00*p))
    dmd('## Negative Categories')
    for k, v, p in zip(list_classes, y[sent_idx], c_pred):
        if v<1:
            dmd('- {}, Prediction: {:.2f}%'.format(k, 100*p))
show_sentence(0)
show_sentence(50)

```

## Input Sentence:

Sure, that couldn't hurt. I've actually tested it on Mac/Firefox and Mac/Safari (albeit an earlier version of Firefox) and it's worked for me. So I think it's less likely a straight-up browser/OS issue, and more likely an interaction with either the monobook.js or perhaps another piece of Javascript code that may be running on your client. But that's just a guess. If it behave differently for you with Safari, please let me know. '''''' Talk

## Positive Categories

## Negative Categories

- toxic, Prediction: 24.37%
- severe\_toxic, Prediction: 23.33%
- obscene, Prediction: 23.79%
- threat, Prediction: 23.26%
- insult, Prediction: 23.75%
- identity\_hate, Prediction: 23.32%

## Input Sentence:

I don't believe it is an attack, but a reminder not to make

Positive Categories

Negative Categories

- toxic, Prediction: 24.37%
- severe\_toxic, Prediction: 23.33%
- obscene, Prediction: 23.79%
- threat, Prediction: 23.26%
- insult, Prediction: 23.75%
- identity\_hate, Prediction: 23.32%

Make Predictions

Load the model and make predictions on the test dataset

In [9]:

```
model.load_weights(file_path)
y_test = model.predict(X_te, verbose = True, batch_size = 1024)
sample_submission = pd.read_csv("../input/sample_submission.csv")
sample_submission[list_classes] = y_test
sample_submission.to_csv("predictions.csv", index=False)
```

This kernel has been released under the [Apache 2.0](#) open source license.

Did you find this Kernel useful?  
Show your appreciation with an upvote

1

153164/153164 [=====] - 92s 603us/step

Data

Data Sources

▼ 🏆 Toxic Comment Clas...

Toxic Comment Classification Challenge



samp...

153k x 7

test.c...

153k x 2

test\_l...

153k x 7

train....

160k x 8

Identify and classify toxic online comments

Last Updated: a year ago

About this Competition

You are provided with a large number of Wikipedia comments which have been labeled by human raters for toxic behavior. The types of toxicity are:

- toxic
- severe\_toxic
- obscene
- threat
- insult
- identity\_hate

You must create a model which predicts a probability of each type of toxicity for each comment.

File descriptions

- train.csv** - the training set, contains comments with their binary labels
- test.csv** - the test set, you must predict the toxicity probabilities for these comments. To deter hand labeling, the test set contains some comments which are not included in scoring.
- sample\_submission.csv** - a sample submission file in the correct format
- test\_labels.csv** - labels for the test data; value of -1 indicates it was not used for scoring; (**Note:** file added after competition close!)

Output Files

New Kernel

Download All

Output Files

predictions.csv

weights.hdf5

About this file

This file was created from a Kernel, it does not have a description.

predictions.csv

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
1							
2	00001cee341fdb12	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
3	0000247867823ef7	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187

4	00013b17ad220c46	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
5	00017563c3f7919a	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
6	00017695ad8997eb	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
7	0001ea8717f6de06	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
8	00024115d4cbde0f	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
9	000247e83dcc1211	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
10	00025358d4737918	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
11	00026d1092fe71cc	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
12	0002eadc3b301559	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
13	0002f87b16116a7f	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
14	0003806b11932181	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
15	0003e1ccfd5a40a	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
16	00059ace3e3e9a53	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
17	000634272d0d44eb	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
18	000663aff0fffc80	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187
19	000689dd34e20979	0.24373632669448853	0.2333270013332367	0.23787954449653625	0.2325991988182068	0.2374608963727951	0.233263242187

## Run Info

Succeeded	True	Run Time	2940.2 seconds
Exit Code	0	Queue Time	0 seconds
Docker Image Name	/python(Dockerfile)	Output Size	0
Timeout Exceeded	False	Used All Space	False
Failure Message			

## Log

Download

```

Time   Line #  Log Message
3.4s   1       [NbConvertApp] Converting notebook script.ipynb to html
3.4s   2       [NbConvertApp] Executing notebook with kernel: python3
44.8s  3       2018-07-12 22:51:42.838822: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:898]
successful NUMA node read from SysFS had negative value (-1),
but there must be at least one NUMA node, so returning NUMA node
zero
44.9s  4       2018-07-12 22:51:42.839549: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1355] Found
device 0 with properties:
name: Tesla K80 major: 3 minor: 7 memoryClockRate(GHz): 0.8235

```

```
pciBusID: 0000:00:04.0
totalMemory: 11.17GiB freeMemory: 11.10GiB
2018-07-12 22:51:42.839611: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1434] Adding
visible gpu devices: 0
47.3s      5 2018-07-12 22:51:45.350638: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:922] Device
interconnect StreamExecutor with strength 1 edge matrix:
2018-07-12 22:51:45.350695: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:928]      0
2018-07-12 22:51:45.350707: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:941] 0:  N
47.4s      6 2018-07-12 22:51:45.351066: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1052] Created
TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0
with 10629 MB memory) -> physical GPU (device: 0, name: Tesla
K80, pci bus id: 0000:00:04.0, compute capability: 3.7)
2939.7s    7 [NbConvertApp] Writing 345963 bytes to __results__.html
2939.7s    8
2939.7s   10 Complete. Exited with code 0.
```


Comments (2)

Sort by

All Comments ▾

Hotness

Please [sign in](#) to leave a comment.




Kevin Mader

Kernel Author

• Posted on Version 1 • 9 months ago

^ 0

@datacanary this early enough to get any Kaggle swag?



DataCanary

• Posted on Latest Version • 9 months ago

^ 0 ▾

Nice! I'll be in touch directly. Incidentally sorry about the slow reply, I didn't get a notification from this at-mention (I've filed a bug report).

Similar Kernels

